

RNN INTRO

+ TEXT SPECIFICALLY

martin@reddragon.ai

sam@reddragon.ai

16 October 2017

WiFi : SG-Guest

Problems with Installation? **ASK!**

PLAN OF ACTION

TODAY

- Start with CNN features wrap-up
- Text & Embeddings
- RNN Intro + Text
- Char-RNN
- RNN Project

PLAN OF ACTION

WEDNESDAY

- Public holiday
 - == Work on projects at home
- Actual session deferred to 'after' :
 - Expected : Monday 20-Nov-2017

PLAN OF ACTION

THURSDAY (OPTIONAL)

- PyTorch Deep Learning MeetUp Group
 - Hosted by Facebook
- Typical Contents :
 - ~~Talk for people starting out~~
 - Things from the bleeding-edge
 - Other talks (including 'lightning')
- [MeetUp.com / PyTorch-and-Deep-Learning-Singapore](https://www.meetup.com/PyTorch-and-Deep-Learning-Singapore/)

PLAN OF ACTION

MONDAY (!)

- RNN project wrap-up
- 'Advanced CNNs'

OUTLINE

- Recurrent Neural Networks
 - Basic idea & problems...
 - GRUs and LSTMs
- Natural Language Processing
 - Tokenization, etc
 - Word Embeddings
- Application : UPPER-CASE NER

PROCESSING SEQUENCES

- Methods so far have fixed input size
- But that is untrue in many domains :
 - Text data (sequences of words)
 - Text data (sequences of characters)
 - Dialog (eg: chatbots)
 - Audio speech data
 - Video clips / movies
- Need a technique that can be applied iteratively

PROCESSING SEQUENCES


VARIABLE-LENGTH INPUT DOESN'T "FIT"

- Run network for each timestep
 - ... with the same parameters
- But 'pass along' internal state
- This state is 'hidden depth'
 - ... and should learn features that are useful

RECURRENT NEURAL NETWORKS

- Apply one network at each step of the input
- Including an internal state that carries forward step-wise
- Everything is still *differentiable*

BASIC RNN

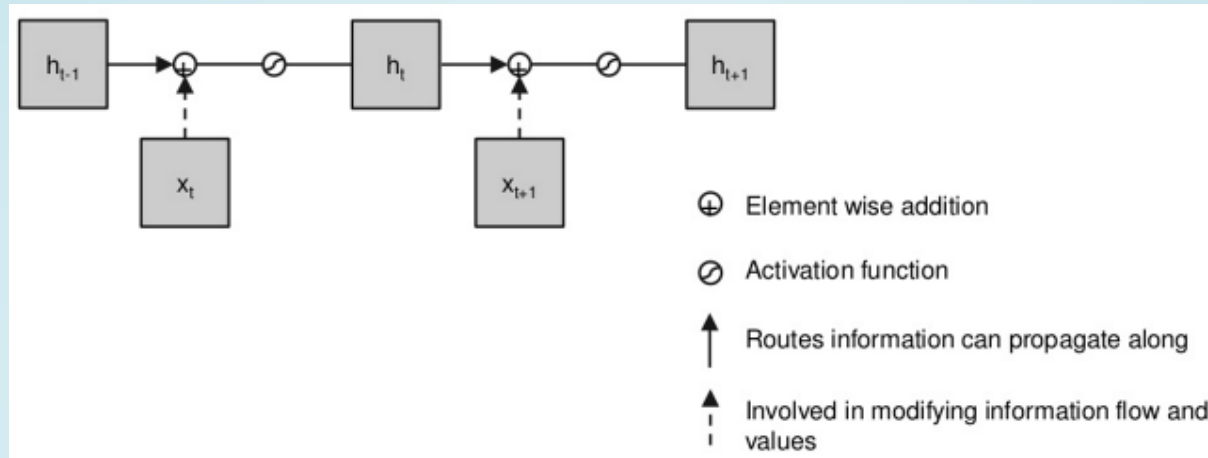


RNN chain

CHAINING STATE

- Each node 'knows' history
- ... all weights are 'tied'
- Network depth is time-wise

PLAIN RNN



Simplest RNN (but has gradient problem)

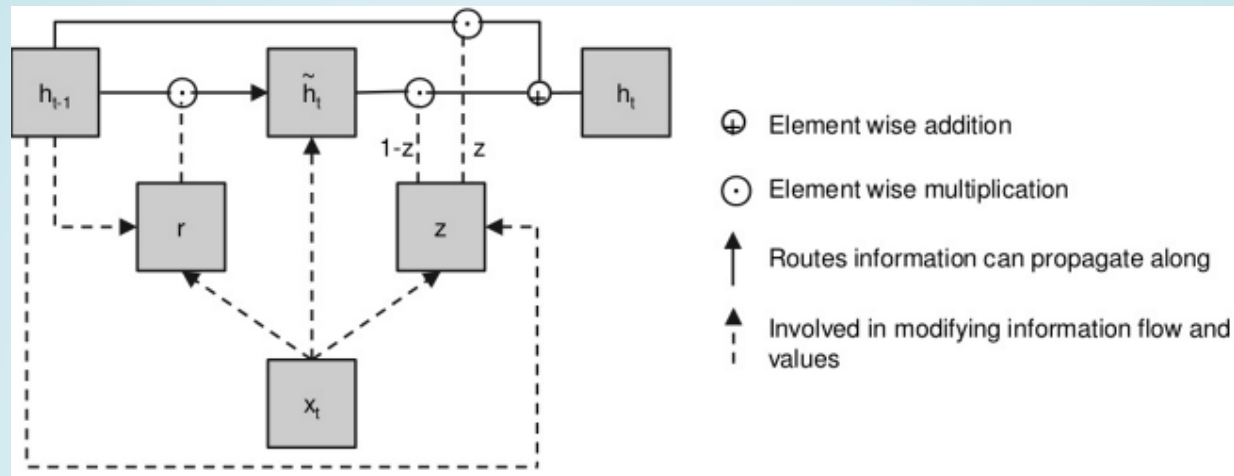
GRADIENT PROBLEM

- In a long sequence, early inputs are 'deep'
- Propagating errors must go through many layers
- Each layer multiplies by an additional factor
- So gradients can explode or vanish

SOLUTION TO GRADIENT PROBLEM

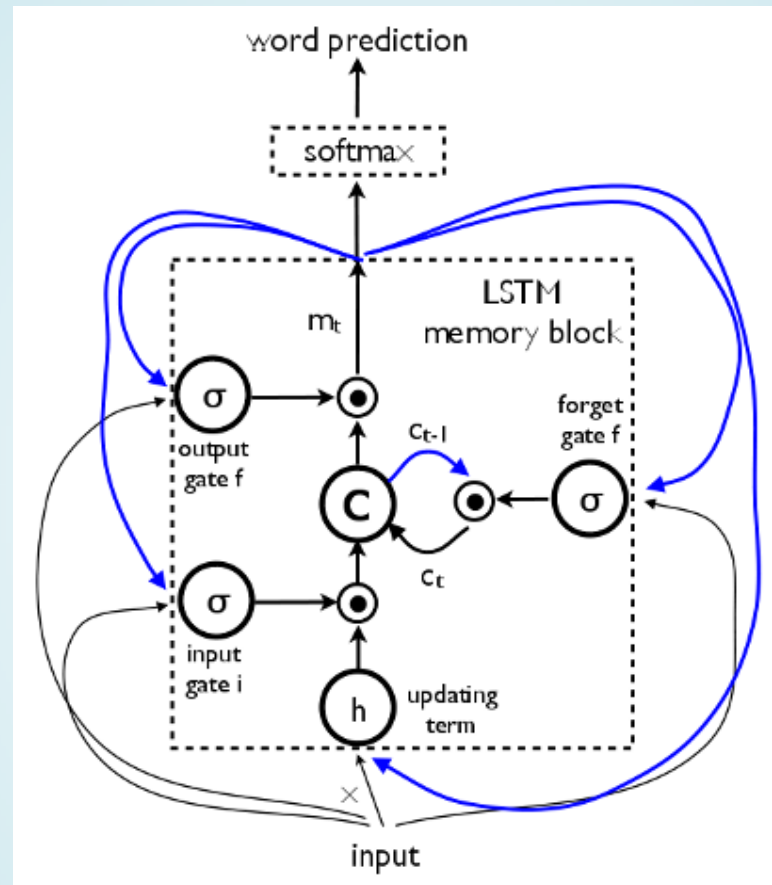
- Instead of always multiplying by weights ...
 - ... have a straight-through path
 - ... so there's a 'gradient 1.0' route
- Idea : Switch the routing on and off

GATED RECURRENT UNITS



A GRU

LSTM UNITS



A Long Short-Term Memory (LSTM) Unit

PILED HIGHER AND DEEPER

- Can also pile up layers
- ... and run forwards and backwards

KEY POINT

- Everything is still *differentiable*
- So : RNNs can be trained
- So : RNNs might be able to learn to do NLP

TEXT

- Documents
- Paragraphs
- Sentences
- Words
- Characters

PRE-PROCESSING

- Encoding
- Sentence-splitting
- Tokenization
- Vocabulary
- ... exceptions

ENCODING

- Can you open the file ?
- Are you 'unicode clean' ?
- Is 'e' the same as 'é' ?
- Is '»' the same as '"' or '"' ?
- Is '.' the same as '•' ?
- How about 'fi' vs 'fi' ?

SENTENCE SPLITTING

Mr. Sam Smith earned \$2012.00
in the U.S.A. in 2012.

TOKENISATION

- Good idea to have a single standard
- Popular approach : Penn-TreeBank (PTB)
- Need to think again for :
 - Chinese (no word spaces)
 - Japanese (unicode punctuation)
 - ... (research is English-centric)
- This is n't so easy .

VOCABULARY BUILDING

- Suppose we have converted sentences to tokens
- Build a dictionary and convert tokens to indices
- Simple frequency analysis :
 - Stop-words (very frequent)
 - Common / normal words
 - Very rare words
 - Typos and Junk
 - <UNK>

UNDERSTANDING TEXT

- English ~100k+ words
- Bag-of-words
- Word Embeddings

BAG-OF-WORDS

- Just convert a sentence into set of words
 - (Throw away ordering)
- Simple statistical analysis (TF-IDF) :
 - Often very effective
- No idea that "jumps" ~ "jump"
- No idea that "jump" ~ "spring"
- No idea that "spring" ~ "summer"

WORD EMBEDDINGS

- Major advances : word2vec & GloVe
- Words that are close in the text should have close representations
- Assign a vector (~300d) to each word
 - Slide a 'window' over the text
 - Vectors for words in window are nudged together
- Keep iterating until 'good enough'
- (corpus size : 1 billion words +)


WORD EMBEDDING

- A map from "token" → Float[100]
- Train over corpus on windows of words
- The vector-space of words self-organizes...

FRANCE	JESUS	XBOX	REDDISH	SCRATCHED	MEGABITS
454	1973	6909	11724	29869	87025
AUSTRIA	GOD	AMIGA	GREENISH	NAILED	OCTETS
BELGIUM	SATI	PLAYSTATION	BLUISH	SMASHED	MB/s
GERMANY	CHRIST	MSX	PINKISH	PUNCHED	BIT/s
ITALY	SATAN	IPOD	PURPLISH	POPPED	BAUD
GREECE	KALI	SEGA	BROWNISH	CRIMPED	CARATS
SWEDEN	INDRA	PSNUMBER	GREYISH	SCRAPED	KBIT/s
NORWAY	VISHNU	HD	GRAYISH	SCREWED	MEGAHERTZ
EUROPE	ANANDA	DREAMCAST	WHITISH	SECTIONED	MEGAPIXELS
HUNGARY	PARVATI	GEFORCE	SILVERY	SLASHED	GBIT/s
SWITZERLAND	GRACE	CAPCOM	YELLOWISH	RIPPED	AMPERES

(eg: word2vec or GloVe)

EMBEDDING VISUALISATION



Highlighting Analogies

TEXT DEMO TIME

[github.com/mdda/deep-learning-workshop/
notebooks/5-RNN/
5-Text-Corpus-and-Embeddings.ipynb](https://github.com/mdda/deep-learning-workshop/notebooks/5-RNN/5-Text-Corpus-and-Embeddings.ipynb)

APPLICATION : RNNS FOR TEXT

- Building a quality NLP system
- Essential component : Named Entity Recognition (NER)
- Has to be flexible / trainable
- Has to understand regional quirks

NER: QUICK EXAMPLE

- Transform :
 - Soon after his graduation , Jim Soon became Managing Director of Lam Soon .
- Into :
 - Soon after his graduation , Jim_Soon^{PER} became Managing Director of Lam_Soon^{ORG} .

LEARNING NAMED ENTITY RECOGNITION

- Can we train an RNN to do 'NER'?
- Steps :
 - Create / select a word embedding
 - Get a NER annotated dataset
 - Train an RNN on the dataset


DEMO DATASET

- Human annotated Corpora are difficult to distribute :
 - Use NLTK to annotate Wikipedia
 - Train RNN on machine annotations
 - Look at performance vs NLTK

LET'S MAKE IT "INTERESTING"

- **Twist : Restrict RNN input to single case text**

NETWORK PICTURE



Bidirectional GRU RNN

RNN-TAGGER-KERAS

[github.com/mdda/deep-learning-workshop/
notebooks/5-RNN/
6-RNN-Tagger-keras.ipynb](https://github.com/mdda/deep-learning-workshop/notebooks/5-RNN/6-RNN-Tagger-keras.ipynb)

WRAP-UP

- Text processing is messy
- Word Embeddings are magic
- RNNs can be applied to lots of things
- Having a GPU is VERY helpful

- QUESTIONS -

MARTIN.ANDREWS @
REDDRAGON.AI

My blog : <http://mda.net/>

GitHub : [mda](#)